Private Benefits and Social Costs of Artificial Intelligence in Education

Matthias Benitez

Introduction: Ethical Concerns Surrounding Artificial Intelligence in Education

Artificial Intelligence has rapidly embedded in educational settings, offering tools that can enhance productivity, accessibility, and learning efficiency. At the same time, widespread AI use has raised significant ethical concerns regarding academic integrity, fairness, and the long-term value of education. While individual students may benefit from AI through time savings, improved writing quality or assistance with complex material, these private benefits coexist with broader social harms that affect educational institutions, employers, and society as a whole.

The central ethical issue is whether it is fair for individuals to capture the benefits of AI use while imposing costs on others who do not consent to or directly participate in those decisions. These costs include diminished skill development and workforce readiness among graduates, erosion of academic trust, increased monitoring and enforcement expenses, and the gradual dilution of educational credentials as reliable signals of skill and effort. When these harms are dispersed across institutions and future stakeholders, they raise questions of responsibility, justice, and proportionality.

Concerns about fairness are further complicated by unequal access to AI tools. Students with greater financial resources or technical familiarity are often better positioned to use AI effectively, while others may lack access or avoid AI use altogether due to ethical concerns. As a result, honest or low-resource students may face competitive disadvantages, bearing the costs of an environment shaped by widespread AI use without sharing equally in its benefits. This dynamic raises ethical questions about whether current approaches unfairly burden certain groups while rewarding others.

Finally, the ethical challenge is not simply whether AI should be permitted or prohibited, but how educational institutions should respond in a way that balances integrity, opportunity, and learning outcomes. Should responsibility be enforced primarily through punishment and surveillance, or should institutions also recognize and encourage responsible AI use that supports genuine learning. Addressing these questions requires careful consideration of both ethical values, such as fairness, responsibility, and equal opportunity, and the economic incentives that shape individual behavior.

Economic Incentives and Social Costs of AI use

Artificial intelligence functions as a productivity enhancing tool that significantly reduces the cost of accessing, organizing, and synthesizing information. For students, AI offers clear private benefits: it can support late night research, clarify difficult concepts, assist with brainstorming and outlining, and help users engage more efficiently with material that has already been taught. In many respects, AI represents an extension of existing technologies such as search engines and online databases, lowering the friction associated with learning rather than replacing it outright. Given these advantages, the rapid adoption of AI across educational settings is neither surprising nor inherently problematic.

At the individual level, the market price of AI use reflects these private considerations. Students weigh modest monetary costs, such as subscription fees, against substantial gains in efficiency, convenience, and learning support. For many users, particularly those who aim to use AI responsibly, the private benefits clearly outweigh the private costs. Importantly, this private calculation does not signal obvious harm; instead, AI use often appears to be a rational and welfare improving choice from the perspective of the individual student.

However, while the benefits of AI use are largely internalized by individual users, several important social costs are not reflected in the market price of AI tools. One significant unpriced cost arises from unequal access. Students with greater financial resources, stronger technical familiarity, or institutional support are better positioned to use AI effectively, while others may lack access or avoid AI altogether due to ethical concerns or uncertainty about acceptable use. As AI becomes more embedded in coursework, these disparities can place honest or low resource students at a competitive disadvantage, even when they have not chosen to engage in AI assisted work. In this way, the benefits of AI use are unevenly distributed, while the resulting competitive pressures are shared more broadly.

A second major social cost emerges through the institutional response to improper AI use. As AI adoption has expanded, educational institutions have increasingly invested in detection software, monitoring systems, and enforcement mechanisms designed to deter misuse. These tools are expensive to acquire and administer, yet their costs are not borne by individual students who engage in unethical AI use. Instead, enforcement costs are spread across institutions and, indirectly, across all students through tuition, fees, or reduced educational resources. Even students who use AI responsibly, or not at all, are required to bear these additional costs, raising ethical concerns about fairness and proportionality.

Beyond Issues of access and enforcement, widespread AI use also generates broader spillover effects related to skill development, credentialing, and trust. When AI substitutes for independent work in ways that weaken skill development, the resulting costs extend beyond the individual student. Employers may face higher training and screening costs, while educational credentials gradually lose their value as reliable signals of competence. Even limited misuse can contribute to the erosion of confidence in academic outcomes, imposing costs on students and employers who rely on these signals but do not benefit from the behavior that undermines them.

Taken together, these dynamics illustrate a clear divergence between private incentives and social outcomes. While AI use may be privately efficient and often educationally beneficial, its broader costs, particularly those associated with unequal access, institutional enforcement, and credential dilution, are not incorporated into the market price faced by individual users. These unpriced social harms constitute negative externalities, as the private benefits of AI use are captured by individuals while the resulting costs are borne by institutions, employers, and society at large. From a Pigouvian perspective, this gap between private and social costs suggests that relying solely on individual decision making may lead to outcomes that are efficient for individuals but ethically and socially problematic at the collective level (Pigou, 1920).

Policy Approaches to Addressing AI Related Pollution

Once the social costs associated with AI use in education are recognized, the relevant question becomes how institutions should respond. Economic theory offers several policy approaches for addressing activities that generate negative externalities, each relaying on different mechanisms to influence behavior and internalize social costs. In practice, most current institutional responses to AI related concerns rely heavily on command-and-control instruments, supplemented by limited and largely punitive incentive structures.

Command-and-control approaches seek to regulate behavior directly through rules, monitoring, and sanctions. In educational settings, this includes mandatory AI bans, zero tolerance policies, disclosure requirements, and the use of AI detection software to identify improper use. These measures operate by increasing the expected cost of misuse through inspections and penalties, thereby discouraging behavior deemed harmful. From a Pigouvian Perspective, such polices function as implicit taxes on unethical AI use by raising the private cost of behavior that imposes social harm (Pigou, 1920).

While command-and-control instruments may deter certain forms of misuse, they face significant limitations when applied to AI. First, enforcement is costly. Detection software, administrative review, and disciplinary processes require substantial institutional resources, the costs of which are not borne by individual users who misuse AI but are instead distributed across institutions and students more broadly. Second, these approaches face information problems. As AI tools evolve rapidly and increasingly integrated into legitimate learning activities, it becomes difficult to distinguish responsible from irresponsible use with precision. This raises the risk of false positives, uneven enforcement, and an erosion of trust between students and institutions.

More importantly, command-and-control approaches address only one side of the incentive problem. They create strong incentives to avoid using AI improperly, but they provide no corresponding incentives to encourage transparent, ethical, and learning oriented use. Students who engage responsibly receive no benefit beyond avoiding punishment, while those who misuse AI face penalties only if detected. As a result, current polices emphasize deterrence rather than incentive alignment, limiting their effectiveness as a comprehensive response to AI related social costs.

An alternative framework often discussed in environmental economics is Coasean bargaining, which emphasizes voluntary negotiation among affected parties as a means of internalizing externalities (Coase, 1960). In theory, clearly defined rights and low transaction costs can allow parties to reach mutually beneficial agreements without the need for regulation. In the context of AI use in education, however, Coasean solutions face substantial obstacles. Although the boundaries of acceptable use are in most cases clearly defined, the number of affected parties is large, and transaction costs are high. Students, faculty, employers, and institutions all have overlapping and sometimes conflicting interests, making decentralized bargaining impractical in this setting.

John Hasnas's critique of standard public policy analysis further emphasizes the importance of comparing institutional alternatives rather than assuming regulation is the default solution (Hasnas, 1995). Hasnas argues that markets, social norms, and legal processes may sometimes resolve conflicts more effectively than legislation, particularly over time. However, AI adoption presents a rapidly evolving and time sensitive challenge. The pace of technological change makes it unlikely that informal norms or case-by-case resolutions will develop quickly enough to prevent the accumulation of social costs. In such contexts, waiting for decentralized solutions to emerge risks allowing inequities and institutional burdens to intensify.

Collectively, these considerations suggest that neither command-and-control regulation nor decentralized market processes are sufficient on their own. While existing policies implicitly "tax" improper AI use through monitoring and sanctions, they neglect the complementary role of the positive incentives. From a Pigouvian perspective, effective policy should aim not only to discourage harmful behavior but also to encourage socially beneficial outcomes (Pigou, 1920). The absence of subsidies or rewards for ethical, transparent, and learning oriented AI use represents a critical gap in current policy approaches, one that raises important questions about fairness, efficiency, and institutional responsibility.

Competing Ethical Perspectives on AI Regulation

Debates surrounding AI used in education often center on how to balance academic integrity with fairness and opportunity. One ethical perspective emphasizes strong regulatory action to prevent harm. From this view, improper AI use undermines the integrity of education, weakens skill development, and erodes trust in academic credentials. Strict enforcement, monitoring, and sanctions are seen as ethically justified to protect honest students, preserve the value of degrees, and prevent long term harm to employers and society. This perspective prioritizes responsibility and precaution, arguing that allowing misuse risks normalizing behavior that degrades educational outcomes.

An opposing perspective cautions that overly punitive or restrictive approaches may impose unfair burdens. Strict bans and zero tolerance policies can disproportionately affect honest or low resource students, particularly when access to AI tools is unequal or institutional rules are unclear. Heavy reliance on surveillance and punishment may also create adversarial learning environments, reduce autonomy, and discourage experimentation with tools that could otherwise enhance understanding. From this standpoint, ethical concern centers on proportionality, equity, and the risk of treating all AI use as inherently harmful rather than context dependent.

Evaluation of Ethical Tradeoffs and Preferred Approach

While both perspectives raise legitimate concerns, the latter offers a more compelling ethical framework when evaluated alongside economic reasoning. Exclusive reliance on punishment addresses symptoms rather than incentives, deterring misuse without encouraging responsible behavior. Moreover, when enforcement costs and access disparities are considered, strict regulatory approaches risk shifting burdens onto students and institutions that are not responsible for the underlying problem.

A Pigouvian framework better accommodates these tradeoffs by focusing on incentive alignment rather than prohibition. Rather than asking whether AI should be permitted or banned, this approach asks how private decision making can be guided toward socially beneficial outcomes. Policies that rely solely on deterrence fail to internalize social costs effectively because they do not reward behaviors that enhance learning, transparency, and skill development. From an ethical standpoint, a system that combines responsibility with opportunity better reflects principles of fairness, efficiency, and proportionality.

Toward a Fair and Balanced institutional Response

A fair and just institutional response to AI use in education should therefore move beyond command-and-control enforcement toward a more balanced, inventive based approach. While monitoring and sanctions may remain necessary to address clear misuse, they should be complemented by mechanisms that actively encourage ethical and educationally valuable AI use. This could include clearer guidelines, institutional access to approved AI tools, structured assignments that integrate AI transparently, and recognition or credit for responsible use.

By pairing deterrence with positive incentives, institutions can reduce the social costs associated with unequal access, enforcement burdens, and credential dilution while preserving the legitimate benefits AI offers to learning. Such an approach aligns with Pigouvian logic by addressing both sides of the externality problem, discouraging harmful behavior while promoting outcomes that enhance collective welfare. Ultimately, this balanced strategy better supports academic integrity, equity, and long-term educational value than polices focused on solely restriction or punishment.

References

Pigou, A.C (1920). *The economics of welfare.* Macmillian

Coase, R.H. (1960). The problem of social cost. *Journal of Law and Economics*, 3, 1-44

Hasnas, J. (n.d) Two theories of environmental regulation [youtube Lecture]. Environmental Economics 4355 West Texas A&M University